

Get started with Azure Databricks

3 minutes

To use Azure Databricks, you must create an Azure Databricks workspace in your Azure subscription. A workspace is an Azure Databricks deployment in a cloud service account. It provides a unified environment for working with Azure Databricks assets for a specified set of users.

You can create an Azure Databricks workspace by:

- Using the Azure portal user interface.
- Using an Azure Resource Manager (ARM), Bicep or Terraform template.
- Using the New-AzDatabricksWorkspace Azure PowerShell cmdlet.
- Using the az databricks workspace create Azure command line interface (CLI) command.

When you create a workspace, you must specify:

- A **workspace name**.
- Select an available **region**. For available regions, see [Azure services available by region](#).
- A **pricing tier**:
 - **Premium** - Role-based access controls, Unity Catalog, SQL, Mosaic AI, serverless compute, Genie Code, and other enterprise-level features.
 - **Trial** - A 14-day free trial of a Premium-level workspace

ⓘ Note

The Standard tier is no longer available for new workspaces as of April 1, 2026. All new workspaces are created on the Premium tier.

- **Workspace type:**
 - **Serverless** - Pre-configured with serverless compute and managed storage.
Recommended for most use cases.
 - **Hybrid** (also called Classic) - Provisions compute and storage in your Azure subscription.
Best for custom networking or on-premises connectivity.
- **Managed Resource Group name** (optional): an automatically created resource group where Azure provisions and manages the infrastructure resources needed for your Databricks workspace.

💡 Tip

A **Free Edition** is available for students and educators with no time limit and no payment required, though it has daily usage limits and no access to classic compute. See [Sign up for Azure Databricks](#) for a comparison.

The screenshot shows the 'Create an Azure Databricks workspace' wizard in the Azure portal, specifically the 'Basics' tab. The interface includes a top navigation bar with the Microsoft Azure logo, a search bar, and user profile information. The main content area is titled 'Create an Azure Databricks workspace' and features a progress bar with tabs for Basics, Networking, Encryption, Security & compliance, Tags, and Review + create. Under the 'Project Details' section, there is a prompt to 'Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.' Below this, there are two dropdown menus: 'Subscription' (set to 'Subscription') and 'Resource group' (set to '(New) rg-databricks'). A 'Create new' link is visible below the resource group dropdown. The 'Instance Details' section contains four fields: 'Workspace name' (set to 'databricks-ws' with a green checkmark), 'Region' (set to 'West Europe'), 'Pricing Tier' (set to 'Premium (+ Role-based access controls)'), and 'Managed Resource Group name' (with a placeholder 'Enter name for managed resource group'). A blue information box below the pricing tier states: 'We selected the recommended pricing tier for your workspace. You can change the tier based on your needs.' At the bottom, there are three buttons: 'Review + create' (highlighted in blue), '< Previous', and 'Next: Networking >'.

If you decide to create an Azure Databricks deployment using the Azure CLI, this would be the [az databricks workspace](#) command to remember:


```
az databricks workspace create
  --resource-group myresourcegroup \
  --name mydatabricksws \
  --location westus2 \
  --sku premium
```

The equivalent [New-AzDatabricksWorkspace](#) PowerShell cmdlet:

```
New-AzDatabricksWorkspace -Name mydatabricksws -ResourceGroupName myresource-
group -Location westus2 -ManagedResourceGroupName databricks-group -Sku
premium
```

Navigating the Azure Databricks Workspace UI

After you provision an Azure Databricks workspace, you can use the workspace UI to work with data and compute resources. The workspace UI is a web-based user interface where you can create and manage workspace resources, such as Spark clusters, and use notebooks and queries to work with data in files and tables.



The homepage provides shortcuts to common tasks and workspace objects to help you get started. You can import data, create a notebook, create a query and configure an AutoML experiment.

The sidebar shows common Databricks categories (Workspace, Recents, Catalog, Jobs & Pipelines, Compute, Marketplace). It then breaks out by product area:

- **SQL:** SQL Editor, Queries, Dashboards, Genie, Alerts, Query History, SQL Warehouses
- **Data Engineering:** Job Runs, Data Ingestion
- **Machine Learning:** Playground, Experiments, Features, Models, Serving

Select + **New** to:


- **Create workspace objects** such as notebooks, queries, repos, dashboards, alerts, jobs, pipelines, experiments, models, and serving endpoints.
- **Create compute resources** such as clusters, SQL warehouses, and ML endpoints.

Use the top bar to **search** for workspace objects such as notebooks, queries, dashboards, alerts, files, folders, libraries, tables registered in Unity Catalog, jobs, and repos in a single place. You can also access recently viewed objects in the search bar.

The workspace is available in **multiple languages**. To change the workspace language, select your username in the top navigation bar, select **Settings** and go to the **Preferences** tab.

Get help from Genie Code

Genie Code (formerly Databricks Assistant) is an AI-powered pair programmer and support tool that helps you work more efficiently in Databricks by generating, explaining, and fixing code or queries directly in notebooks, dashboards, and files.



It can assist with a wide range of tasks, including identifying and correcting errors, creating data visualizations, diagnosing job issues, and filtering or analyzing data using natural language prompts. Genie Code can surface relevant guidance from the Azure Databricks documentation.

Genie Code also includes **Agent mode**, which is generally available for data science, data engineering, and dashboard authoring. Agent mode extends Genie Code with agentic capabilities so it can autonomously plan and complete multi-step data tasks on your behalf.

By using Unity Catalog metadata, it personalizes its responses based on your organization's data assets—tables, columns, and descriptions—making it easier to explore and work with your data.

All units complete:

[< Previous](#)[Complete module](#)