

Explore Azure Databricks — Knowledge Check

20 questions · Explore Azure Databricks (Microsoft Learn, Unit 1). Answer key starts on a new page.

1. What two technologies power the Databricks lakehouse?
 - A. Hadoop and Hive
 - B. Delta Lake and Unity Catalog
 - C. Spark and DBFS
 - D. Parquet and Kafka
2. What is the "medallion architecture"?
 - A. A pricing tier
 - B. A layered pattern that progressively refines data as it moves through stages
 - C. A type of SQL warehouse
 - D. A Unity Catalog permission model
3. True or False: As of April 1, 2026, new Azure Databricks workspaces can no longer be created on the Standard tier.

4. Which workspace type is recommended for most use cases, offering pre-configured serverless compute and managed storage?
 - A. Hybrid / Classic
 - B. Serverless
 - C. Trial
 - D. Standard
5. Which Azure CLI command creates a new Azure Databricks workspace?

6. What is Genie Code (formerly known as)?
 - A. Databricks Assistant
 - B. Databricks Runtime
 - C. Lakeflow Jobs
 - D. Databricks SQL
7. Which Azure Databricks workload area relies on MLflow for experiment tracking and model management?
 - A. Data Engineering
 - B. Machine Learning
 - C. SQL / BI
 - D. Streaming
8. True or False: SQL warehouses in Azure Databricks require the Premium tier.

9. In the Azure Databricks two-layer architecture, which layer is managed by Microsoft and handles backend services?
 - A. Compute plane
 - B. Control plane

- C. Storage plane
- D. Metastore plane

10. What are the two node types that make up a Databricks cluster?

11. Which Databricks Runtime lifecycle stage means a version is retired and no longer available at all?

- A. Legacy
- B. Deprecated
- C. End of Support (EoS)
- D. End of Life (EoL)

12. What is a Lakeflow Job, structurally?

- A. A single Spark cluster
- B. A container for one or more tasks
- C. A Unity Catalog schema
- D. A SQL warehouse type

13. Which Delta Lake feature lets you query or roll back to a previous version of a table?

- A. Predictive IO
- B. Time travel
- C. Intelligent Workload Management
- D. Photon

14. Which SQL warehouse type has compute running inside your own Azure subscription and takes about 4 minutes to start?

- A. Serverless only
- B. Pro and Classic
- C. Lakehouse Real-Time
- D. None — all warehouse types start instantly

15. Which SQL warehouse type supports Intelligent Workload Management (IWM)?

- A. Classic only
- B. Pro only
- C. Serverless only
- D. All three types equally

16. What is the top-level metadata container in Unity Catalog, typically one per region?

- A. Catalog
- B. Schema
- C. Metastore
- D. Workspace

17. What is Unity Catalog's three-level namespace format?

18. What's the key difference between a managed table and an external table in Unity Catalog?

- A. Managed tables can't be queried with SQL
- B. Unity Catalog controls both governance and storage/lifecycle for managed tables; for external tables it governs access only

- C. External tables are always faster
- D. There is no difference

19. What does Microsoft Purview do when integrated with Unity Catalog?

- A. Replaces Unity Catalog entirely
- B. Discovers Lakehouse data and ingests its metadata into the Purview Data Map for cross-platform governance
- C. Only manages billing
- D. Converts external tables into managed tables

20. True or False: Microsoft Purview can only scan Unity Catalog metadata — it cannot scan the legacy workspace-level Hive metastore.

Answer Key

1. B. Delta Lake and Unity Catalog

Delta Lake provides ACID transactions and schema enforcement; Unity Catalog provides unified governance for data and AI.

2. B. A layered pattern that progressively refines data as it moves through stages

It's a common lakehouse design pattern for gradually cleaning and enriching data.

3. True

Standard stopped accepting new workspaces on April 1, 2026; existing Standard workspaces auto-upgrade to Premium by October 1, 2026.

4. B. Serverless

Hybrid (Classic) is best reserved for custom networking or on-premises connectivity needs instead.

5. az databricks workspace create

The equivalent PowerShell cmdlet is New-AzDatabricksWorkspace.

6. A. Databricks Assistant

Genie Code is the AI pair programmer built into notebooks, dashboards, and files; its Agent mode can autonomously complete multi-step tasks.

7. B. Machine Learning

MLflow manages the end-to-end ML lifecycle, including reproducibility and deployment, and also supports GenAI workflows.

8. True

Databricks SQL and SQL warehouses are Premium-tier-only capabilities.

9. B. Control plane

The compute plane is the layer that actually processes data and lives in your own Azure subscription.

10. Driver node and worker node(s)

The driver node coordinates execution; worker nodes handle the distributed computation. Clusters can be fixed-size or auto-scaling.

11. D. End of Life (EoL)

The order is: Legacy → Deprecated → End of Support (no more patches) → End of Life (retired, unavailable).

12. B. A container for one or more tasks

Tasks define the actual work (running a notebook, a Spark job, external code, etc.); jobs can trigger on a schedule, an event, or manually.

13. B. Time travel

Delta Lake's versioning/rollback capability is commonly called time travel, enabled by its transaction log.

14. B. Pro and Classic

Both Pro and Classic warehouses run compute in your subscription and take roughly 4 minutes to start; Serverless starts in about 2-6 seconds.

15. C. Serverless only

Per the performance capability table: Serverless supports Photon + Predictive IO + IWM; Pro supports Photon + Predictive IO; Classic supports Photon only.

16. C. Metastore

Multiple workspaces in the same region can share a single metastore, giving them a unified view of governed data.

17. catalog.schema.object

Catalogs group assets by team/environment, schemas subdivide by project/use case, and objects include tables, views, volumes, functions, and models.

18. B. Unity Catalog controls both governance and storage/lifecycle for managed tables; for external tables it governs access only

Managed tables are always stored as Delta Lake; external tables can be in other formats (CSV, JSON, Parquet, etc.) with lifecycle managed outside Databricks.

19. B. Discovers Lakehouse data and ingests its metadata into the Purview Data Map for cross-platform governance

Purview can scan the entire Unity Catalog metastore (or selected catalogs), auto-classify sensitive data, and show cross-system lineage.

20. False

Purview can also scan the workspace-level Hive metastore in Azure Databricks, in addition to Unity Catalog.