

Understand Azure Databricks Integrations — Knowledge Check

20 questions · Understand Azure Databricks integrations (Microsoft Learn, Unit 3). Answer key starts on a new page.

1. What governance layer does every Azure Databricks integration in this module preserve, even though enforcement mechanisms differ?

- A. Azure RBAC
- B. Unity Catalog
- C. Microsoft Entra Conditional Access
- D. DBFS permissions

2. When Fabric reads Unity Catalog data through a Mirrored Azure Databricks Catalog, whose identity is used for authorization?

- A. Each individual querying user's identity
- B. The identity of whoever configured the connection
- C. A random service account
- D. No authorization is performed

3. What Unity Catalog privilege must a user hold to set up Fabric's read access to a schema?

4. True or False: Databricks can write data to a Fabric lakehouse in OneLake using Entra credential passthrough on serverless compute.

5. Since February 2026, what driver do new Power BI connections to Azure Databricks use by default?

- A. Simba Spark ODBC Driver
- B. JDBC
- C. ADBC (Arrow Database Connectivity)
- D. Thrift

6. What does publishing data from Azure Databricks to the Power BI service require?

- A. Nothing beyond a free Power BI account
- B. Unity Catalog and a Power BI Premium license with the XMLA endpoint enabled
- C. A classic SQL warehouse only
- D. Government cloud access

7. In the Power BI authentication comparison, which method allows Unity Catalog to enforce row-level filters and column masks per individual viewer?

- A. Personal access token
- B. Import mode with no SSO
- C. Microsoft Entra ID (SSO)
- D. Service principal

8. What VS Code capability lets you set breakpoints and step through code that's actually executing on a remote Databricks cluster?

- A. Declarative Automation Bundles
- B. Databricks Connect
- C. Partner Connect
- D. One-way sync

9. What were Databricks Asset Bundles renamed to?

10. True or False: The VS Code extension's workspace sync feature is two-way — edits made directly in the Databricks workspace sync back to your local project.

11. What compute resource does the Azure Databricks connector for Power Platform require?

- A. A classic cluster
- B. A SQL warehouse
- C. Serverless notebooks only
- D. DBFS mounts

12. Which Power Apps data integration option does NOT support OAuth credential passthrough?

- A. Canvas apps with a direct connection
- B. Dataverse virtual tables
- C. Power Automate SQL execution
- D. Job orchestration

13. True or False: The Azure Databricks Power Platform connector supports US Government and China Cloud environments.

14. In Copilot Studio, what's the key difference between using tables as 'knowledge sources' versus connecting a 'Genie space' as a tool?

- A. There is no difference
- B. Knowledge sources do semantic search/retrieval; Genie performs full analytics, aggregation, and trend explanation
- C. Genie spaces don't use Unity Catalog governance
- D. Knowledge sources require MCP, Genie doesn't

15. Which SQL warehouse types are supported for Genie in Copilot Studio?

- A. Classic only
- B. Serverless or pro only
- C. Any warehouse type
- D. None — Genie doesn't use SQL warehouses

16. How does Microsoft Purview access data when scanning Azure Databricks?

- A. It copies the underlying data into Purview
- B. It reads only metadata — databases, tables, columns, relationships — never the actual data
- C. It requires exporting data to CSV first
- D. It only works with external tables

17. What's a key difference between Hive Metastore scanning and Unity Catalog scanning in Purview?

- A. Hive Metastore scanning supports incremental sync; Unity Catalog doesn't

- B. Unity Catalog scanning supports incremental sync and captures runtime lineage; Hive Metastore scanning does full extraction only with static lineage
- C. They are functionally identical
- D. Hive Metastore scanning requires no infrastructure at all

18. What protocol connects Microsoft Foundry AI agents to Azure Databricks Genie spaces?

19. What is the Genie API rate limit when called from Microsoft Foundry?

- A. 5 questions per minute
- B. 100 questions per minute
- C. Unlimited
- D. 1 question per hour

20. True or False: As of mid-2026, the Genie integrations with Copilot Studio and Microsoft Foundry are generally available (GA) and fully production-ready.

Answer Key

1. B. Unity Catalog

Unity Catalog stays central across all seven integrations; what changes is how strictly each service enforces its permissions.

2. B. The identity of whoever configured the connection

This is a key security gotcha: Unity Catalog governance doesn't automatically apply to downstream Fabric users unless you also enable OneLake security on the mirrored catalog.

3. EXTERNAL USE SCHEMA

This privilege lets Unity Catalog issue credentials for the underlying storage; it's only needed by the person configuring the connection, not by downstream Fabric users.

4. False

Credential passthrough is not supported with serverless compute — serverless writes to OneLake require service principal authentication instead.

5. C. ADBC (Arrow Database Connectivity)

ADBC is a modern, Arrow-based columnar API; existing connections stay on ODBC unless manually switched, and full transition to ADBC is planned for August 2026.

6. B. Unity Catalog and a Power BI Premium license with the XMLA endpoint enabled

The XMLA endpoint (Read/Write capability) is what lets Databricks programmatically create and update semantic models in the Power BI workspace.

7. C. Microsoft Entra ID (SSO)

Entra ID SSO passes each viewer's own credentials through to Azure Databricks so Unity Catalog evaluates permissions per user.

8. B. Databricks Connect

Databricks Connect powers full interactive debugging in VS Code, including notebooks debugged cell by cell.

9. Declarative Automation Bundles

They package code, configuration, and dependencies (in YAML) into deployable units for Lakeflow Jobs, Spark Declarative Pipelines, and MLOps Stacks.

10. False

Sync is one-way, local to workspace only. Workspace files are meant to be transient; edits made there won't sync back.

11. B. A SQL warehouse

SQL warehouses provide the optimized query performance and serverless scaling the connector relies on.

12. B. Dataverse virtual tables

Virtual tables integrate Databricks data into model-driven apps without copying it, but they trade away OAuth passthrough — direct connections perform better and do support it.

13. False

Government clouds are explicitly unsupported for this connector; organizations in those environments need alternative integration approaches.

14. B. Knowledge sources do semantic search/retrieval; Genie performs full analytics, aggregation, and trend explanation

Genie spaces connect via Model Context Protocol (MCP) and provide a much deeper analytical engine than simple knowledge-source retrieval.

15. B. Serverless or pro only

Classic warehouses are explicitly not supported for Genie integrations.

16. B. It reads only metadata — databases, tables, columns, relationships — never the actual data

The Purview integration works entirely through metadata synchronization: registration, then scanning at scheduled or on-demand intervals.

17. B. Unity Catalog scanning supports incremental sync and captures runtime lineage; Hive Metastore scanning does full extraction only with static lineage

Unity Catalog scanning also uses a cloud-native Azure integration runtime via SQL Warehouses, while Hive Metastore scanning needs a self-hosted integration runtime.

18. Model Context Protocol (MCP)

Foundry acts as an MCP client; the Databricks MCP server exposes Genie spaces (and other resources) as MCP-compatible tools.

19. A. 5 questions per minute

Design agent prompts and UX to avoid rapid sequential queries against this limit; the Genie MCP tool also has no conversation memory, so each question is processed independently.

20. False

Both integrations are in Public Preview (workspace-wide MCP endpoint went live June 2, 2026) — suitable for development/testing, but subject to change before GA.